

Prefix Tuning / P-Tuning v2 实战

为了不影响阅读体验，详细的代码放置在GitHub: llm-action 项目中 [peft_prefix_tuning_clm.ipynb](#) 和 [peft_p_tuning_v2_clm.ipynb](#) 文件，这里仅列出关键步骤。

第一步，引进必要的库，如：Prefix Tuning / P-Tuning v2 配置类 PrefixTuningConfig。

```
1 from peft import get_peft_config, get_peft_model, PrefixTuningConfig, TaskType, PeftType
```

第二步，创建 Prefix Tuning / P-Tuning v2 微调方法对应的配置。

```
1 peft_config = PrefixTuningConfig(task_type=TaskType.CAUSAL_LM, num_virtual_tokens=30)
```

PrefixTuningConfig 配置类参数说明：

- task_type: 指定任务类型。如：条件生成任务 (SEQ_2_SEQ_LM)，因果语言建模 (CAUSAL_LM) 等。
- num_virtual_tokens: 虚拟token的数量，换句话说就是提示 (prompt)。
- inference_mode: 是否在推理模式下使用Peft模型。
- prefix_projection: 是否投影前缀嵌入(token)，默认值为false，表示使用P-Tuning v2， 如果为true，则表示使用 Prefix Tuning。

第三步，通过调用 get_peft_model 方法包装基础的 Transformer 模型。

```
1 model = AutoModelForCausalLM.from_pretrained(model_name_or_path)
2 model = get_peft_model(model, peft_config)
3 model.print_trainable_parameters()
```

通过 print_trainable_parameters 方法可以查看到 P-Tuning v2 可训练参数的数量(仅为1,474,560)以及占比 (仅为0.2629%)。

```
1 trainable params: 1,474,560 || all params: 560,689,152 || trainable%: 0.26299064191632515
```

PEFT 中 Prefix Tuning 相关的代码是基于清华开源的[P-tuning-v2](#) 进行的重构；同时，我们可以在 chatglm-6b 和 chatglm2-6b 中看到类似的[代码](#)。PEFT 中源码如下所示。

```
1 class PrefixEncoder(torch.nn.Module):
2     def __init__(self, config):
3         super().__init__()
4         self.prefix_projection = config.prefix_projection
5         token_dim = config.token_dim
6         num_layers = config.num_layers
7         encoder_hidden_size = config.encoder_hidden_size
8         num_virtual_tokens = config.num_virtual_tokens
9         if self.prefix_projection and not config.inference_mode:
10             # Use a two-layer MLP to encode the prefix
11             # 初始化重参数化的编码器
12             self.embedding = torch.nn.Embedding(num_virtual_tokens, token_dim)
13             self.transform = torch.nn.Sequential(
14                 torch.nn.Linear(token_dim, encoder_hidden_size),
15                 torch.nn.Tanh(),
16                 torch.nn.Linear(encoder_hidden_size, num_layers * 2 * token_dim),
17             )
18         else:
19             self.embedding = torch.nn.Embedding(num_virtual_tokens, num_layers * 2 * token_dim)
20
21     def forward(self, prefix: torch.Tensor):
22         if self.prefix_projection:
23             prefix_tokens = self.embedding(prefix)
24             past_key_values = self.transform(prefix_tokens)
25         else:
26             past_key_values = self.embedding(prefix)
27         return past_key_values
```

从上面的源码也可以看到 Prefix Tuning 与 P-Tuning v2 最主要的差别就是是否进行重新参数化编码。

第四步，模型训练的其余部分均无需更改，当模型训练完成之后，保存高效微调部分的模型权重以供模型推理即可。

▼ Plain Text |

```
1 peft_model_id = f"{model_name_or_path}_{peft_config.peft_type}_{peft_config.task_type}"
2 model.save_pretrained(peft_model_id)
```

输出的模型权重文件如下所示：

▼ Plain Text |

```
1 /data/nfs/llm/model/bloomz-560m_PREFIX_TUNING_CAUSAL_LM
2 |— [ 390] adapter_config.json
3 |— [5.6M] adapter_model.bin
4 |— [ 93] README.md
5
6 0 directories, 3 files
```

注意：这里只会保存经过训练的增量 PEFT 权重。其中，adapter_config.json 为 P-Tuning v2 / Prefix Tuning 配置文件；adapter_model.bin 为 P-Tuning v2 / Prefix Tuning 权重文件。

第五步，加载微调后的权重文件进行推理。

```
1 from peft import PeftModel, PeftConfig
2
3 peft_model_id = f"{model_name_or_path}_{peft_config.peft_type}_{peft_config.task_type}"
4 config = PeftConfig.from_pretrained(peft_model_id)
5 # 加载基础模型
6 model = AutoModelForCausalLM.from_pretrained(config.base_model_name_or_path)
7 # 加载PEFT模型
8 model = PeftModel.from_pretrained(model, peft_model_id)
9
10 # 编码
11 inputs = tokenizer(f'{text_column} : {dataset["test"][i]["Tweet text"]} Label : ', return_tensors="pt")
12
13 # 模型推理
14 outputs = model.generate(
15     input_ids=inputs["input_ids"],
16     attention_mask=inputs["attention_mask"],
17     max_new_tokens=10,
18     eos_token_id=3
19 )
20
21 # 解码
22 print(tokenizer.batch_decode(outputs.detach().cpu().numpy(), skip_special_tokens=True))
```

至此，我们完成了 Prefix Tuning / P-Tuning v2 的训练及推理。

```

1  # coding=utf-8
2  # Copyright 2023-present the HuggingFace Inc. team.
3  #
4  # Licensed under the Apache License, Version 2.0 (the "License");
5  # you may not use this file except in compliance with the License.
6  # You may obtain a copy of the License at
7  #
8  #     http://www.apache.org/licenses/LICENSE-2.0
9  #
10 # Unless required by applicable law or agreed to in writing, software
11 # distributed under the License is distributed on an "AS IS" BASIS,
12 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
13 # See the License for the specific language governing permissions and
14 # limitations under the License.
15
16 # Based on https://github.com/THUDM/P-tuning-v2/blob/main/model/prefix_encoder.py
17 # with some refactor
18 import torch
19
20
21 class PrefixEncoder(torch.nn.Module):
22     """
23     The `torch.nn` model to encode the prefix.
24
25     Args:
26         config ([`PrefixTuningConfig`]): The configuration of the prefix encoder.
27
28     Example:
29
30     ```py
31     >>> from peft import PrefixEncoder, PrefixTuningConfig
32
33     >>> config = PrefixTuningConfig(
34         ...     peft_type="PREFIX_TUNING",
35         ...     task_type="SEQ_2_SEQ_LM",
36         ...     num_virtual_tokens=20,
37         ...     token_dim=768,
38         ...     num_transformer_submodules=1,
39         ...     num_attention_heads=12,
40         ...     num_layers=12,
41         ...     encoder_hidden_size=768,
42         ... )
43     >>> prefix_encoder = PrefixEncoder(config)

```

```

44     ...
45
46     **Attributes**:
47         - **embedding** (`torch.nn.Embedding`) -- The embedding layer of t
48 he prefix encoder.
49         - **transform** (`torch.nn.Sequential`) -- The two-layer MLP to tr
50 ansform the prefix embeddings if
51         `prefix_projection` is `True`.
52         - **prefix_projection** (`bool`) -- Whether to project the prefix
53 embeddings.
54
55     Input shape: (`batch_size`, `num_virtual_tokens`)
56
57     Output shape: (`batch_size`, `num_virtual_tokens`, `2*layers*hidden`)
58     """
59
60     def __init__(self, config):
61         super().__init__()
62         self.prefix_projection = config.prefix_projection
63         token_dim = config.token_dim
64         num_layers = config.num_layers
65         encoder_hidden_size = config.encoder_hidden_size
66         num_virtual_tokens = config.num_virtual_tokens
67         if self.prefix_projection and not config.inference_mode:
68             # Use a two-layer MLP to encode the prefix
69             self.embedding = torch.nn.Embedding(num_virtual_tokens, token_
70 dim)
71             self.transform = torch.nn.Sequential(
72                 torch.nn.Linear(token_dim, encoder_hidden_size),
73                 torch.nn.Tanh(),
74                 torch.nn.Linear(encoder_hidden_size, num_layers * 2 * toke
75 n_dim),
76             )
77         else:
78             self.embedding = torch.nn.Embedding(num_virtual_tokens, num_la
79 yers * 2 * token_dim)
80
81     def forward(self, prefix: torch.Tensor):
82         if self.prefix_projection:
83             prefix_tokens = self.embedding(prefix)
84             past_key_values = self.transform(prefix_tokens)
85         else:
86             past_key_values = self.embedding(prefix)
87         return past_key_values

```